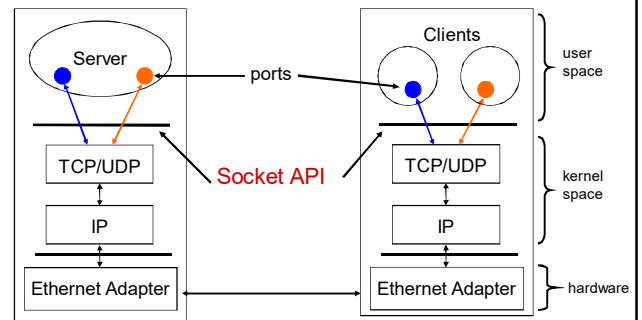


## Socket programming in C


## Server and Client

Server and Client exchange messages over the network through a common **Socket API**



## Concept of Port Numbers

- Port numbers are used to identify “entities” on a host (16 bit)
- Port numbers can be
  - Well-known (port 0-1023)
  - Registered (port 1024 – 49151)
  - Dynamic or private (port 49152-65535)
- Servers/daemons usually use well-known ports
  - Any client can identify the server/service
  - HTTP = 80, FTP = 21, Telnet = 23, ...
  - `/etc/services` defines well-known ports
- Clients usually use dynamic ports
  - Assigned by the kernel at run time



## Names and Addresses

- Each attachment point on Internet is given unique address
  - Based on location within network – like phone numbers
- Humans prefer to deal with names not addresses
  - DNS provides mapping of name to address
  - Name based on administrative ownership of host

## Internet Addressing Data Structure

```
#include <netinet/in.h>

/* Internet address structure */
struct in_addr {
    u_long s_addr; /* 32-bit IPv4 address */
}; /* network byte ordered */

/* Socket address, Internet style. */
struct sockaddr_in {
    u_char sin_family; /* Address Family */
    u_short sin_port; /* UDP or TCP Port# */
    /* network byte ordered */
    struct in_addr sin_addr; /* Internet Address */
    char sin_zero[8]; /* unused */
};
```

- `sin_family = AF_INET` selects Internet address family

## Byte Ordering

```
union {
    u_int32_t addr; /* unsigned long int, 4
                     bytes address */
    unsigned char c[4];
} un;
/* 128.2.194.95 */
un.addr = 0x8002c25f;
/* un.c[0] = ? */
```

- Big Endian → 

128	2	194	95
-----	---	-----	----


  
– Sun Solaris, PowerPC, ...
- Little Endian → 

95	194	2	128
----	-----	---	-----


  
– i386, alpha, ...
- Network byte order = Big Endian

### Byte Order

- Different computers may have different internal representation of 16 / 32-bit integer (called **host byte order**).
- Examples
  - Big-Endian byte order (e.g., used by Motorola 68000):



- Little-Endian byte order (e.g., used by Intel 80x86):



## Byte Ordering Functions

- Converts between **host byte order** and **network byte order**
  - ‘h’ = host byte order
  - ‘n’ = network byte order
  - ‘l’ = long (4 bytes), converts IP addresses
  - ‘s’ = short (2 bytes), converts port numbers

```
#include <netinet/in.h>

unsigned long int htonl(unsigned long int hostlong);
unsigned short int htons(unsigned short int hostshort);
unsigned long int ntohl(unsigned long int netlong);
unsigned short int ntohs(unsigned short int netshort);
```


## What is a Socket?

- A socket is a file descriptor that lets an application read/write data from/to the network

```
int fd;          /* socket descriptor */
if ((fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- `socket` returns an integer (socket descriptor)
  - `fd < 0` indicates that an error occurred
  - socket descriptors are similar to file descriptors
- `AF_INET`: associates a socket with the Internet protocol family
- `SOCK_STREAM`: selects the TCP protocol
- `SOCK_DGRAM`: selects the UDP protocol

## TCP Server



- For example: web server
- What does a *web server* need to do so that a *web client* can connect to it?

## Socket I/O: socket()

- Since web traffic uses TCP, the web server must create a socket of type `SOCK_STREAM`

```
int fd;          /* socket descriptor */

if ((fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- `socket` returns an integer (socket descriptor)
  - `fd < 0` indicates that an error occurred
- `AF_INET` associates a socket with the Internet protocol family
- `SOCK_STREAM` selects the TCP protocol

## Socket I/O: bind()

- A *socket* can be bound to a *port*

```
int fd;          /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */

/* create the socket */

srv.sin_family = AF_INET; /* use the Internet addr family */

srv.sin_port = htons(80); /* bind socket 'fd' to port 80*/

/* bind: a client may connect to any of server addresses */
srv.sin_addr.s_addr = htonl(INADDR_ANY);

if (bind(fd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("bind"); exit(1);
}
```

- Still not quite ready to communicate with a client...

## Socket I/O: listen()

- **listen** indicates that the server will accept a connection

```
int fd; /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */

/* 1) create the socket */
/* 2) bind the socket to a port */

if(listen(fd, 5) < 0) {
    perror("listen");
    exit(1);
}
```

- Still not quite ready to communicate with a client...

## Socket I/O: accept()

- **accept** blocks waiting for a connection

```
int fd; /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */
struct sockaddr_in cli; /* client's add used by accept() */
int newfd; /* returned by accept() */
int cli_len = sizeof(cli); /* used by accept() */

/* 1) create the socket */
/* 2) bind the socket to a port */
/* 3) listen on the socket */

newfd = accept(fd, (struct sockaddr*) &cli, &cli_len);
if(newfd < 0) {
    perror("accept"); exit(1);
}
```

- **accept** returns a new socket (**newfd**) with the same properties as the original socket (**fd**)
  - **newfd** < 0 indicates that an error occurred

## Socket I/O: accept() continued...

```
struct sockaddr_in cli; /* used by accept() */
int newfd; /* returned by accept() */
int cli_len = sizeof(cli); /* used by accept() */

newfd = accept(fd, (struct sockaddr*) &cli, &cli_len);
if(newfd < 0) {
    perror("accept");
    exit(1);
}
```

- How does the server know which client it is?
  - **cli.sin\_addr.s\_addr** contains the client's **IP address**
  - **cli.sin\_port** contains the client's **port number**
- Now the server can exchange data with the client by using **read** and **write** on the descriptor **newfd**.
- Why does **accept** need to return a new descriptor?

## Socket I/O: read()

- **read** can be used with a socket
- **read blocks** waiting for data from the client but does not guarantee that **sizeof(buf)** is read

```
int fd; /* socket descriptor */
char buf[512]; /* used by read() */
int nbytes; /* used by read() */


/* 1) create the socket */
/* 2) bind the socket to a port */
/* 3) listen on the socket */
/* 4) accept the incoming connection */
/* 5) Create a child process to attend the connection */
/* Child will close fd, and continue communication with client */
/* parent will close newfd, and continue with accepting connection*/

if((nbytes = read(newfd, buf, sizeof(buf))) < 0) {
    perror("read"); exit(1);
}
```

## TCP Client

- For example: web client

- **How does a web client connect to a web server?**



## Dealing with IP Addresses

- IP Addresses are commonly written as strings ("128.2.35.50"), but programs deal with IP addresses as integers.

### Converting strings to numerical address:

```
struct sockaddr_in srv;

srv.sin_addr.s_addr = inet_addr("128.2.35.50");
if(srv.sin_addr.s_addr == (in_addr_t) -1) {
    fprintf(stderr, "inet_addr failed!\n"); exit(1);
}
```

### Converting a numerical address to a string:

```
struct sockaddr_in srv;
char *t = inet_ntoa(srv.sin_addr);
if(t == 0) {
    fprintf(stderr, "inet_ntoa failed!\n"); exit(1);
}
```

## Translating Names to Addresses

- Gethostbyname provides interface to DNS
- Additional useful calls
  - Gethostbyaddr – returns hostent given sockaddr\_in
  - Getservbyname
    - Used to get service description (typically port number)
    - Returns servent based on name

```
#include <netdb.h>

struct hostent *hp; /*ptr to host info for remote*/
struct sockaddr_in peeraddr;
char *name = "www.tezu.ernet.in";

peeraddr.sin_family = AF_INET;
hp = gethostbyname(name);
peeraddr.sin_addr.s_addr = ((struct in_addr*) (hp->h_addr))->s_addr;
```

## Socket I/O: connect()

- **connect** allows a client to connect to a server...

```
int fd; /* socket descriptor */
struct sockaddr_in srv; /* used by connect() */

/* create the socket */

/* connect: use the Internet address family */
srv.sin_family = AF_INET;

/* connect: socket 'fd' to port 80 */
srv.sin_port = htons(80);

/* connect: connect to IP Address "128.2.35.50" */
srv.sin_addr.s_addr = inet_addr("128.2.35.50");

if(connect(fd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("connect"); exit(1);
}
```

## Socket I/O: write()

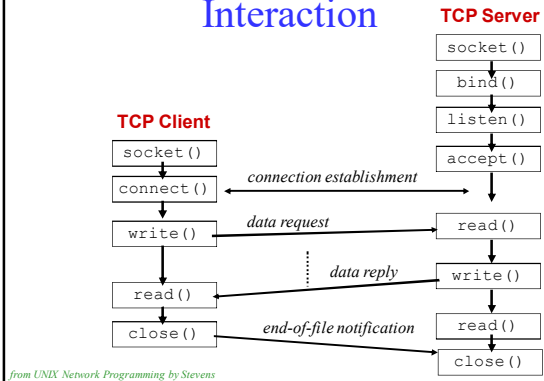
- write** can be used with a socket

```
int fd; /* socket descriptor */
struct sockaddr_in srv; /* used by connect() */
char buf[512]; /* used by write() */
int nbytes; /* used by write() */

/* 1) create the socket */
/* 2) connect() to the server */

/* Example: A client could "write" a request to a server */
if((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

## Review: TCP Client-Server Interaction



## Socket programming *with TCP*

### Client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

### Client contacts server by:

- creating client-local TCP socket
- specifying IP address, port number of server process
- When **client creates socket**: client TCP establishes connection to server TCP

- When contacted by client, **server TCP creates new socket** for server process to communicate with client
  - allows server to talk with multiple clients
  - source port numbers used to distinguish clients.

**application viewpoint**  
*TCP provides reliable, in-order transfer of bytes ("pipe") between client and server*

## Example: C client (TCP)

```
/* client.c */
void main(int argc, char *argv[])
{
    struct sockaddr_in sad; /* structure to hold an IP address */
    int clientSocket; /* socket descriptor */
    struct hostent *ptrh; /* pointer to a host table entry */

    char Sentence[128];
    char modifiedSentence[128];

    host = argv[1]; port = atoi(argv[2]);

    clientSocket = socket(PF_INET, SOCK_STREAM, 0);
    memset((char *)&sad, 0, sizeof(sad)); /* clear sockaddr structure */
    sad.sin_family = AF_INET; /* set family to Internet */
    sad.sin_port = htons((u_short)port);
    ptrh = gethostbyname(host); /* Convert host name to IP address */
    memcpy(&sad.sin_addr, ptrh->h_addr, ptrh->h_length); connect(clientSocket,
    (struct sockaddr *)&sad, sizeof(sad));
}
```

Create client socket, connect to server

## Example: C client (TCP), cont.

```

    Get
input stream  } gets(Sentence);
from user

    Send line  } n=write(clientSocket, Sentence, strlen(Sentence)+1);
to server

    Read line  } n=read(clientSocket, modifiedSentence, sizeof(modifiedSentence));
from server

    printf("FROM SERVER: %s\n",modifiedSentence);

    Close     } close(clientSocket);
connection
}

```

## Example: C server (TCP)

```

/* server.c */
void main(int argc, char *argv[])
{
    struct sockaddr_in sad; /* structure to hold an IP address */
    struct sockaddr_in cad;
    int welcomeSocket, connectionSocket; /* socket descriptor */
    struct hostent *ptrh; /* pointer to a host table entry */

    char clientSentence[128];
    char capitalizedSentence[128];

    port = atoi(argv[1]);

    Create welcoming socket at port
    &
    Bind a local address
    ↓
    welcomeSocket = socket(PF_INET, SOCK_STREAM, 0);
    memset((char *)&sad,0,sizeof(sad)); /* clear sockaddr structure */
    sad.sin_family = AF_INET; /* set family to Internet */
    sad.sin_addr.s_addr = INADDR_ANY; /* set the local IP address */
    sad.sin_port = htons((u_short)port); /* set the port number */
    bind(welcomeSocket, (struct sockaddr *)&sad, sizeof(sad));
}

```

## Example: C server (TCP), cont

```

/* Specify the maximum number of clients that can be queued */
listen(welcomeSocket, 10)

while(1) {
    connectionSocket=accept(welcomeSocket, (struct sockaddr *)&cad, &alen);
    // Wait, on welcoming socket for contact by a client

    n=read(connectionSocket, clientSentence, sizeof(clientSentence));

    /* capitalize Sentence and store the result in capitalizedSentence*/

    n=write(connectionSocket, capitalizedSentence, strlen(capitalizedSentence)+1);
    // Write out the result to socket

    close(connectionSocket);
}
// End of while loop, loop back and wait for another client connection

```

## Socket programming *with UDP*

UDP: no "connection"  
between client and server

- no handshaking
- sender explicitly attaches IP address and port of destination to each packet
- server must extract IP address, port of sender from received packet

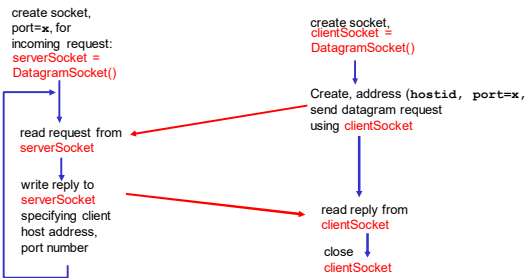
UDP: transmitted data may  
be received out of order,  
or lost

application viewpoint  
UDP provides unreliable transfer  
of groups of bytes ("datagrams")  
between client and server

## Client/server socket interaction: UDP

Server (running on `hostid`)

Client



## Example: C client (UDP)

```

/* client.c */
void main(int argc, char *argv[])
{
    struct sockaddr_in sad; /* structure to hold an IP address */
    int clientSocket; /* socket descriptor */
    struct hostent *ptrh; /* pointer to a host table entry */

    char Sentence[128];
    char modifiedSentence[128];

    host = argv[1]; port = atoi(argv[2]);

    clientSocket = socket(PF_INET, SOCK_DGRAM, 0);

    /* determine the server's address */
    memset((char *)&sad, 0, sizeof(sad)); /* clear sockaddr structure */
    sad.sin_family = AF_INET; /* set family to Internet */
    sad.sin_port = htons((u_short)port);
    ptrh = gethostbyname(host); /* Convert host name to IP address */
    memcpy(&sad.sin_addr, ptrh->h_addr, ptrh->h_length);
  
```

*Create client socket, NO connection to server*

## Example: C client (UDP), cont.

```

    Get input stream from user ]-> gets(Sentence);

    Send line to server ]-> addr_len = sizeof(struct sockaddr);
    n=sendto(clientSocket, Sentence, strlen(Sentence)+1,
             (struct sockaddr *) &sad, addr_len);

    Read line from server ]-> n=recvfrom(clientSocket, modifiedSentence, sizeof(modifiedSentence),
                                         (struct sockaddr *) &sad, &addr_len);

    printf("FROM SERVER: %s\n", modifiedSentence);

    Close connection ]-> close(clientSocket);
  }
  
```

## Example: C server (UDP)

```

/* server.c */
void main(int argc, char *argv[])
{
    struct sockaddr_in sad; /* structure to hold an IP address */
    struct sockaddr_in cad;
    int serverSocket; /* socket descriptor */
    struct hostent *ptrh; /* pointer to a host table entry */

    char clientSentence[128];
    char capitalizedSentence[128];

    port = atoi(argv[1]);

    serverSocket = socket(PF_INET, SOCK_DGRAM, 0);
    memset((char *)&sad, 0, sizeof(sad)); /* clear sockaddr structure */
    sad.sin_family = AF_INET; /* set family to Internet */
    sad.sin_addr.s_addr = INADDR_ANY; /* set the local IP address */
    sad.sin_port = htons((u_short)port); /* set the port number */
    bind(serverSocket, (struct sockaddr *)&sad, sizeof(sad));
  
```

*Create welcoming socket at port & Bind a local address*



## Example: C server (UDP), cont

```
while(1) {
    n=recvfrom(serverSocket, clientSentence, sizeof(clientSentence), 0
               (struct sockaddr *) &cad, &addr_len );

    /* capitalize Sentence and store the result in capitalizedSentence*/

    n=sendto(connectionSocket, capitalizedSentence, strlen(capitalizedSentence)+1,0
              (struct sockaddr *) &cad, &addr_len);

    close(connectionSocket);
}
```

Receive messages from clients

Write out the result to socket

End of while loop, loop back and wait for another client connection

## A Few Programming Notes: Representing Packets

0				1				2				3									
0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1
Type																					
Length										Checksum											
Address																					

Type: 4-byte integer  
 Length: 2-byte integer  
 Checksum: 2-byte integer  
 Address: 4-byte IP address

## A Few Programming Notes: Building a Packet in a Buffer

```
struct packet {
    u_int32_t type;
    u_int16_t length;
    u_int16_t checksum;
    u_int32_t address;
};

/* ===== */
char buf[1024];
struct packet *pkt;

pkt = (struct packet*) buf;
pkt->type = htonl(1);
pkt->length = htons(2);
pkt->checksum = htons(3);
pkt->address = htonl(4);
```

## Socket Programming References

- Man page
  - usage: man <function name>
- Textbook
  - Stevens, Unix Network Programming, PHI